

functions.py

December 17, 2018

```
In [ ]: # Functions for project
```

```
import numpy as np
import pandas as pd
import re
import tensorflow as tf
from sklearn.model_selection import train_test_split
import matplotlib
import matplotlib.pyplot as plt
from mpl_toolkits.axes_grid1 import make_axes_locatable
import random
from sklearn.metrics import r2_score, mean_squared_error, accuracy_score, log_loss
from sklearn import svm #support vector machines
from sklearn.feature_extraction.text import CountVectorizer

from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neural_network import MLPClassifier

from sklearn.metrics import confusion_matrix
```

```
#####
#####
# Heatmap plotting functions
#####
#####
```

```
##### This is code to do neat plotting, taken from:
# https://matplotlib.org/gallery/images_contours_and_fields/image_annotated_heatmap.ht
```

```
def heatmap(data, row_labels, col_labels, ax=None,
            cbar_kw={}, cbarlabel="", **kwargs):
    """
    Create a heatmap from a numpy array and two lists of labels.
```

Arguments:

data : A 2D numpy array of shape (N,M)
row_labels : A list or array of length N with the labels for the rows
col_labels : A list or array of length M with the labels for the columns

Optional arguments:

ax : A matplotlib.axes.Axes instance to which the heatmap is plotted. If not provided, use current axes or create a new one.
cbar_kw : A dictionary with arguments to :meth:`matplotlib.figure.colorbar`.
cbarlabel : The label for the colorbar

All other arguments are directly passed on to the imshow call.
"""

```
if not ax:
    ax = plt.gca()

# Plot the heatmap
im = ax.imshow(data, **kwargs)

# Create colorbar
cbar = ax.figure.colorbar(im, ax=ax, **cbar_kw)
cbar.ax.set_ylabel(cbarlabel, rotation=-90, va="bottom")

# We want to show all ticks...
ax.set_xticks(np.arange(data.shape[1]))
ax.set_yticks(np.arange(data.shape[0]))
# ... and label them with the respective list entries.
ax.set_xticklabels(col_labels)
ax.set_yticklabels(row_labels)

# Let the horizontal axes labeling appear on top.
ax.tick_params(top=True, bottom=False,
               labeltop=True, labelbottom=False)

# Rotate the tick labels and set their alignment.
plt.setp(ax.get_xticklabels(), rotation=-30, ha="right",
         rotation_mode="anchor")

# Turn spines off and create white grid.
for edge, spine in ax.spines.items():
    spine.set_visible(False)

ax.set_xticks(np.arange(data.shape[1]+1)-.5, minor=True)
ax.set_yticks(np.arange(data.shape[0]+1)-.5, minor=True)
```

```

ax.grid(which="minor", color="w", linestyle='-', linewidth=3)
ax.tick_params(which="minor", bottom=False, left=False)

return im, cbar

def annotate_heatmap(im, data=None, valfmt="{x:.2f}",
                    textcolors=["black", "white"],
                    threshold=None, **textkw):
    """
    A function to annotate a heatmap.

    Arguments:
        im : The AxesImage to be labeled.
    Optional arguments:
        data : Data used to annotate. If None, the image's data is used.
        valfmt : The format of the annotations inside the heatmap.
                This should either use the string format method, e.g.
                "$ {x:.2f}", or be a :class:`matplotlib.ticker.Formatter`.
        textcolors : A list or array of two color specifications. The first is
                    used for values below a threshold, the second for those
                    above.
        threshold : Value in data units according to which the colors from
                    textcolors are applied. If None (the default) uses the
                    middle of the colormap as separation.

    Further arguments are passed on to the created text labels.
    """

    if not isinstance(data, (list, np.ndarray)):
        data = im.get_array()

    # Normalize the threshold to the images color range.
    if threshold is not None:
        threshold = im.norm(threshold)
    else:
        threshold = im.norm(data.max())/2.

    # Set default alignment to center, but allow it to be
    # overwritten by textkw.
    kw = dict(horizontalalignment="center",
                verticalalignment="center")
    kw.update(textkw)

    # Get the formatter in case a string is supplied
    if isinstance(valfmt, str):
        valfmt = matplotlib.ticker.StrMethodFormatter(valfmt)

```

```

    # Loop over the data and create a `Text` for each "pixel".
    # Change the text's color depending on the data.
    texts = []
    for i in range(data.shape[0]):
        for j in range(data.shape[1]):
            kw.update(color=textcolors[im.norm(data[i, j]) > threshold])
            text = im.axes.text(j, i, valfmt(data[i, j], None), **kw)
            texts.append(text)

    return texts

#####
#####
# plotting code over
#####
#####

#####
#####
# plotting code for confusion matrices
#####
#####

def plot_confusion_matrix(prediction, true_vals, labels, size = (12,12), normalize = 'rows'):
    confusion = confusion_matrix(prediction, true_vals, labels = labels)
    fig, ax = plt.subplots(figsize=size)

    if normalize == 'rows' :
        im, cbar = heatmap((confusion.T/confusion.sum(axis=1)).T, labels, labels, ax = ax, cmap = "YlGn")
    elif normalize == 'columns' :
        im, cbar = heatmap(confusion/confusion.sum(axis=1), labels, labels, ax = ax, cmap = "YlGn")
    else :
        im, cbar = heatmap(confusion, labels, labels, ax = ax, cmap = "YlGn", cbarlabel = None)

    texts = annotate_heatmap(im, valfmt="{x:.2f}")
    fig.tight_layout()
    plt.show()

def clf_confusion(clf, x_train, y_train, x_test, y_test, labels, size = (12,12), normalize = 'rows'):
    clf.fit(x_train, y_train)
    predictions = clf.predict(x_test)
    plot_confusion_matrix(predictions, y_test, labels, size = size, normalize = normalize)

#####
#####

```

```

# plotting code for confusion matrices over
#####
#####

#####
#####
# Functions for reading in the data
#####
#####

def get_design_matrix(cleaning_function = lambda x : x, min_df = 0.0, max_df = 1.0) :
    """
    Take a data frame data, and convert to a matrix.
    Use cleaning_function to clear up data.
    """
    data = pd.read_json('train.json')
    recipie_list_list = data.ingredients.values.tolist()
    recipie_string_list = [cleaning_function(" ".join(ing)) for ing in recipie_list_list]
    vectorizer = CountVectorizer(min_df = min_df, max_df = max_df)
    X = vectorizer.fit_transform(recipie_string_list)
    y = data.cuisine.values
    return X, y, vectorizer.get_feature_names()

def clean(s) :
    clean_s = s.replace('-', ' ') # treat low-fat and low fat as the same thing
    clean_s = ''.join([c for c in clean_s if (c.isalpha() or c == ' ')]) # drop numbers
    return clean_s

#####
#####
# Method validation functions
#####
#####

def clf_cross_validator(X_train, y_train, clf_constructor, p_list, q_list = [], folds = 10) :
    """
    A general method to preform cross validation of sci_kit learn method.
    Takes:
    Training data (X_train, y_train)
    A function clf_constructor which builds a sci_kit learn classifier
    p_list a list of parameters to varry
    q_list a possible second list to varry,
    folds the number of folds to use for cross validation
    """

```

```

plot determines if a heatmap of the results should be printed
label is the label of the plotting
"""

scores = []
# we rename the constructor,
# only plays a role if q_list is empty
constructor = clf_constructor

####
# If we are only passed one list,
# modify the constructor to take a second dummy argument
# We make q_list be a singleton list,
# and flip the role of p_list and q_list, the latter being only for printing purpo
if not q_list : # true if q_list is empty
    q_list = p_list
    p_list = ['']
    constructor = (lambda p,q : clf_constructor(q))

###
# We loop over the parameters
# and record the avarage of each folds score
for p in p_list :
    print("    p=%s" % str(p))
    for q in q_list :
        print("        q=%s" % str(q))
        clf = constructor(p,q)
        score = cross_val_score(clf, X_train, y_train, cv=folds)
        scores.append(np.mean(score))

# transform the scores to a len(p_list) x len(q_list) shape array
scores_array = np.array(scores).reshape(len(p_list), len(q_list))

####
# make a heat map of the scores
if plot :
    fig, ax = plt.subplots()
    im, cbar = heatmap(scores_array, np.array(p_list), np.array(q_list) , ax = ax,
        texts = annotate_heatmap(im, valfmt="{x:.4f}")
    fig.tight_layout()
    plt.show()

return scores_array

def svm_tester(X_train, y_train, C_list = [0.1], folds = 10, plot = False) :
    """
    Test the svm parameter C using cross validation

```

```

For each C in C_list do folds fold cross validation,
returns an array, of the same size as C_list,
of the avarages of the accuracies for each fold.

If plot is set to true, show a heatmap of the results
"""

svm_constructor = (lambda p : svm.LinearSVC(C = p))
scores = clf_cross_validator(X_train, y_train, svm_constructor, C_list, folds = fo
return scores

def forrest_tester(X_train, y_train, trees_list = [1], depth_list = [1], folds = 10, p
"""
Test the random forrest for the parameters of number of trees and max depth using
For each pair of parameters do folds fold cross validation,
returns an array, of shape len(trees_list) x len(depth_list)
of the avarages of the accuracies for each fold.

If plot is set to true, show a heatmap of the results
"""

forrest_constructor = (lambda p,q : RandomForestClassifier(n_estimators = p, max_de
scores = clf_cross_validator(X_train, y_train, forrest_constructor, trees_list, de
return scores

def logistic_tester(X_train, y_train, C_list = [0.1], folds = 10, plot = False) :
"""
Test the logistic regression parameter C using cross validation
For each C in C_list do folds fold cross validation,
returns an array, of the same size as C_list,
of the avarages of the accuracies for each fold.

If plot is set to true, show a heatmap of the results
"""

# Not really sure about the solver
logistic_constructor = (lambda p : LogisticRegression(solver='lbfgs', multi_class=
scores = clf_cross_validator(X_train, y_train, logistic_constructor, C_list, folds
return scores

def mlp_tester(X_train, y_train, nodes = [1], alpha_list = [0.001], folds = 10, plot =
"""
Test the mlp classifier (neural net) for the parameters of
number of nodes in a single layer and regularization constant using cross validati
For each pair of parameters do folds fold cross validation,
returns an array, of shape len(nodes) x len(alpha_list)
of the avarages of the accuracies for each fold.

```

If plot is set to true, show a heatmap of the results
 """

```
mlp_constructor = (lambda p,q : MLPClassifier(hidden_layer_sizes = p, alpha = q, max_iter=1000))
scores = clf_cross_validator(X_train, y_train, mlp_constructor, nodes, alpha_list,
return scores
```

```
def accuracy_with_min_df(min_df_list = [0], svm_parms = [0.1], log_parms = [1], forrest_parms = [10]):
    svm_accuracies = []
    log_accuracies = []
    forrest_accuracies = []

    # rather than calling get_design_matrix every time,
    # which loads the file train.json every time,
    # we just read it in once.
    data = pd.read_json('train.json')
    recipie_list_list = data.ingredients.values.tolist()
    recipie_string_list = [" ".join(ing) for ing in recipie_list_list]
    y = data.cuisine.values
    del data, recipie_list_list

    for df in min_df_list :
        print("Testing min_df = %f" % df)
        vectorizer = CountVectorizer(min_df = df)
        X = vectorizer.fit_transform(recipie_string_list)
        # Pick the highest cv scores (average of fold accuracies)
        print("Now doing svm cross validation")
        svm_accuracies.append(np.amax(svm_tester(X, y, C_list = svm_parms, folds = 10)))
        print("Now doing logistic cross validation")
        log_accuracies.append(np.amax(logistic_tester(X, y, C_list = log_parms, folds = 10)))
        print("Now doing forrest cross validation")
        forrest_accuracies.append(np.amax(forrest_tester(X, y, trees_list = forrest_parms, folds = 10)))
        print("")

    return svm_accuracies, log_accuracies, forrest_accuracies
```